Chapter 6

The EM algorithm and Information Theory

6.1 Mixture Models

In the previous chapter we have mentioned that it may happen that a likelihood function has multiple maxima and that sometimes it may be hard to impossible to find the global maximum (i.e. the maximum with the overall highest likelihood value). Such a situation occurs whenever the probabilistic model that we use to model our observations is a **mixture model**.

Definition 6.1 (Mixture Model) *Given any set of probability distributions* $P_{X_1}, \ldots, P_{X_n}$ *over the same RV X we define a mixture model as* $P = \sum_{i=1}^{n} \alpha_i P_{X_i}$ *where we require that all* $\alpha_i \geq 0$ *and* $\sum_{i=1}^{n} \alpha_i = 1$. *We call the distributions* $P_{X_1}, \ldots, P_{X_n}$ ***mixture components*** *and their weights* α_i ***mixture weights***.

Mixture models are extremely useful whenever we have different ways to think about our data. Each way of conceptualising our data can be encoded by one of the mixture components of the mixture model. This point of view can help us to build a better overall model of our data. Let us introduce a running example that we use for the rest of this section.

Example of a mixture model Assume we observe 20 sequences of coin tosses. Each sequence contains 10 flips. We also know that there are 3 coins with which these sequences could possibly be generated and for each sequence a different coin may have been used. Coin 1 is unbiased, coin 2 has parameter $\theta = 0.4$ and coin 3 has parameter $\theta = 0.65$.

We could assume that the entire data set has been generated by exactly one coin. We would then employ maximum likelihood estimation and pick the coin to find the parameter of that coin. However, this model might

actually turn out to be pretty bad because we are committing to picking only one coin, although we said in the beginning that each sequence may possibly have been generated by a different coin.

A mixture model comes to the rescue. Instead of assuming that only one coin has generated all 20 sequences, we assume that all three coins have contributed to generating the 20 sequences. However, their contributions may not be equal. This inequality is exactly what the mixture weights capture. As usual, we call our data x. Also, each mixture component is a binomial distribution, parametrized by the parameters of the coins. These observations suffice to formulate our mixture model.

$$
\begin{aligned}
P(X = x|\Theta_1^3 = \theta_1^3, A = \alpha_1^3) = {}& \alpha_1 P(X = x|\Theta_1 = 0.4) \\
& + \alpha_2 P(X = x|\Theta_2 = 0.5) + \alpha_3 P(X = x|\Theta_3 = 0.65)
\end{aligned}
$$
(6.1)

Notice that if we were given the mixture weights, estimating the parameters of the mixture components would be easy: we would simply find the MLE for each mixture component. The mixture model could then easily be constructed because the mixture weights are known.

Usually, we face the opposite problem with mixture models. We know the mixture components, but do not know the mixture weights. How can we go about estimating the weights? First, observe the constraints on mixture weights in Definition 6.1. All mixture weights have to be non-negative and they have to sum to 1. These constraints mean that we can interpret the weights as a probability distribution. In particular, the weights form a distribution over parameters of the mixture components (assuming that all mixture components have the same parametric form). That is, we get the following relations:

$$
(6.2) \qquad \alpha_1 = P(\Theta_1 = \theta_1) \quad \alpha_2 = P(\Theta_2 = \theta_1) \quad \alpha_3 = P(\Theta_3 = \theta_1)
$$

Hence, we can rewrite our mixture model as

(6.3)
$$
\begin{aligned}
P(X = x) = {}& P(\Theta_1 = 0.4)P(X = x|\Theta_1 = 0.4) \\
& + P(\Theta_2 = 0.5)P(X = x|\Theta_2 = 0.5) + P(\Theta_3 = 0.65)P(X = x|\Theta_3 = 0.65)
\end{aligned}
$$

Notice that in our examples we assume that we have only three possible binomial parameters, namely $\theta_1 = 0.4, \theta_2 = 0.5, \theta_3 = 0.65$. Notice further that each summand in (6.3) is in fact a joint distribution by the chain rule so that we can again rewrite the model as

$$
\begin{aligned}
P(X = x) = {}& P(X = x, \Theta_1 = 0.4) + P(X = x, \Theta_2 = 0.5) \\
& + P(X = x, \Theta_3 = 0.65) \\
= {}& \sum_{i=1}^{3} P(X = x, \Theta = \theta_i)
\end{aligned}
$$
(6.4)

where (6.4) is a simple marginalisation step. Notice that we have just accomplished something impressive: we have given a probabilistic justification for why mixture models do indeed model our data. Each mixture-component/weight pair gives rise to a joint distribution over parameters and data but by summing over the possible parameters we get the probability of the data. We can easily show that mixture models can be defined for any number of mixture components.

Exercise 6.2 *Show that a mixture model of size n is a model of the data, i.e. show that $\sum_{i=1}^{n} \alpha_i P_{X_i}(X = x) = P(X = x)$ if α_i are mixture weights as defined in Definition 6.1.*

Notice that we have shown above that the formulation of mixture models as purely probabilistic model and as linear combinations of probability distributions are equivalent. So why did we even bother to give a purely probabilistic justification? On the one hand, it is mathematically satisfying to trace back new concepts to concepts that we are already familiar with (like joint distributions and the chain rule). More importantly, however, the probabilistic interpretation of mixture weights allows to estimate them using the maximum-likelihood principle. This estimation is something we could not have done, if we had regarded them solely as scale factors for the mixture components.

There is one additional problem, however: if the mixture weights are unknown, there is no closed-form solution for estimating the likelihood function. This is because the likelihood depends on the parameters of the mixture components their priors, the mixture weights, are unknown. In other words, we can not simply apply calculus as we have been doing up to now. For this reason, we turn to the **EM algorithm**, that allows to at least find a local maximum of the likelihood function.

6.2   The EM algorithm

In order to estimate the parameters of mixture models, we can employ a classical algorithm of **unsupervised learning**, namely the **expectation-maximisation (EM) algorithm**. This algorithm allows to find a local maximum of the likelihood function of mixture models or, more generally, models with missing data.

Let us quickly introduce the idea of missing data: Assume you run a website that recommends movies based on a user's preferences. In order to make statistical predictions about what type of user likes what kind of movie, you ask your users to rate movies according to different categories. Say you ask your users to rate the movies for entertainment value, action and fun. What may happen is that some of your users only rate a movie in one or two of these three categories. However, these ratings are still valuable to

you and you do not want to throw them away, just because they are lacking a rating in one category. Thus you have a data set with some missing data that you have to fill in somehow.

In mixture models, annotations that tell you which mixture component generated a given data point can be thought of as missing data. If you had this information, you could simply do maximum likelihood estimation. However, since you do not know which mixture component generated your data point, you can only assume that all mixture components may have done so with a probability that is given by the posterior distribution over mixture components.

The EM algorithm allows to probabilistically fill in the missing data and find good mixture weights (where you should understand *good* in the maximum-likelihood sense). The idea behind the algorithm is simple: compute the expected number of occurrences of the missing data values (the mixture components) and then do maximum likelihood estimation on those expectations. Repeat the procedure until the likelihood does not increase any further. Notice that this procedure requires to fix the number of mixture components in advance.

More formally, assume a data set x_1^n. Furthermore, define Y as a random variable indicating m possible mixture components. In the more general case of missing data, Y is called **latent** or **hidden variable** because it is not observed. Then the likelihood function is

$$(6.5) \qquad L_x(\theta) = P(X = x | \Theta = \theta) = \sum_{y=1}^{m} P(X = x, Y = y | \Theta = \theta)$$

where Θ ranges over the parameters of the joint distribution P_{XY}. Let $t(y)$ be the statistic that allows for the computation of the MLE. For all distributions that we have seen so far and in fact for all distributions in the exponential family, this statistic is the number of times we observe a particular outcome y. The formulation that we give of the EM algorithm is specific for exponential family distributions and can be made more specific. However, since virtually all distributions that are of interest in practice do belong to the exponential family, we will not make this kind of generalisation.

Given some estimate $\theta^{(i)}$ for the parameters of the joint distribution compute the expected value of the statistic $t(y)$.

$$(6.6) \qquad\qquad t(y)^{(i+1)} = \mathbb{E}(t(y) | X = x, \Theta = \theta^{(i)})$$

Observe that we introduced superscripts. The EM algorithm is an iterative algorithm, meaning we repeat its steps several times. We use the the superscript to indicate the number i of the repetition, thus $0 \le i \le m$. We pass on information from previous iterations by conditioning on their parameter estimates; $\theta^{(0)}$ can be set arbitrarily.

Equation (6.6) is know as the **E(xpectation)-step** of the EM algorithm. To make the algorithm complete, we are still lacking a **M(aximization)-step**. But that step is simple. We pretend that the expected statistics of the latent data were actually observed. If our statistics are counts, then we simply pretend that we observed outcome y exactly $t(y)^{(i+1)}$ times. Notice that since $t(y)^{(i+1)}$ is an expectation, it is a real and not a natural number. Once we pretend to observe the expected statistics, the maximization step can be done using be maximum-likelihood estimation:

$$(6.7) \qquad \theta^{(i+1)} = \arg\max_\theta P(X = x, t(Y) = t(y)|\theta)$$

Definition 6.3 (EM algorithm) *We assume a data set $x = x_1^n$ and postulate that there is unobserved data $y = y_1^m$ that belongs to that data set. We also assume a probabilistic model whose parameters are realisations of a RV Θ. Let $t(y)$ be the count of outcome y. Then any iterative algorithm with k iterations that performs the following steps for $0 \leq i \leq k - 1$,*

> ***E-step:*** $\quad t(y)^{(i+1)} = \mathbb{E}(t(y) \mid X = x, \Theta = \theta^{(i)})$
>
> ***M-step:*** $\quad \theta^{(i+1)} = \arg\max_\theta P(X = x, t(Y) = t(y)^{(i+1)} \mid \Theta = \theta)$

to update the model parameters is called an EM algorithm.

Example of an EM algorithm   Assume as in Section 6.1 that our data is $x = x_1^{10}$ where each x_i is the number of heads that we observed in a sequence of a hundred coin tosses. Again we also assume mixture components that are binomials with parameters $0.4, 0.5, 0.65$. The latent data in this case is an annotation that for each observed sequence x_i reveals the coin that has been used to generate that sequence. Thus we have latent data $y = y_1^n$ with $y_i \in \text{supp}(Y) = \{0.4, 0.5, 0.65\}$. Both the observed and latent variables are assumed i.i.d. We assume that the fair coin is more likely to be used and hence set its initial mixture weight to 0.5 and the mixture weights of the other two coins to 0.25 (any other choice would also be fine). Let us take a closer look at our data. To shorten notation, we write it as a list where the ith entry is the value of x_i.

$$[6, 5, 4, 2, 6, 6, 6, 5, 4, 2, 5, 5, 3, 4, 6, 4, 5, 6, 3, 3]$$

Then for each x_i we assume that it was generated by each of the three coins. For the first observation we get the following likelihood values.

$$(6.8) \qquad P(X_1 = 6|\Theta = 0.4) = 0.1114767$$
$$P(X_1 = 6|\Theta = 0.5) = 0.2050781$$
$$P(X_1 = 6|\Theta = 0.65) = 0.2376685$$

Recall that the mixture weights are nothing else than priors over mixture components. Hence, in order to get the joint distribution over observed and latent data, we multiply the likelihoods by the mixture weights.

$$(6.9) \quad P(X_1 = 6, \Theta = 0.4) = 0.25 \times P(X_1 = 6|\Theta = 0.4) = 0.02786918$$
$$P(X_1 = 6, \Theta = 0.5) = 0.5 \times P(X_1 = 6|\Theta = 0.4) = 0.1025391$$
$$P(X_1 = 6, \Theta = 0.65) = 0.25 \times P(X_1 = 6|\Theta = 0.4) = 0.05941712$$

We are interested in the expected number of times that each coin has generated x_1. Since x_1 has been generated by exactly one coin, the highest expected count is 1. To see this, simply imagine that each sequence of tosses is padded with one more number indicating the coin used to produce the sequence. Then clearly, this coin is observed once as a sequence generator for every sequence that it is padded to.

In order to compute an expectation we first need a distribution over the mixture components. This distribution is simply the posterior. The posterior given x_1 is shown below.

$$(6.10) \quad P(\Theta = 0.4|X_1 = 6) = \frac{P(X_1 = 6, \Theta = 0.4)}{P(X_1 = 6)} = 0.146814$$

$$P(\Theta = 0.5|X_1 = 6) = \frac{P(X_1 = 6, \Theta = 0.5)}{P(X_1 = 6)} = 0.5401758$$

$$P(\Theta = 0.65|X_1 = 6) = \frac{P(X_1 = 6, \Theta = 0.65)}{P(X_1 = 6)} = 0.3130094$$

We can compute the expectation. As statistic for the latent value $a \in \text{supp}(Y)$ **Chris:** correct? **Philip:** In principle correct. I just used a instead of y here because I wanted to remind people what an indicator function is in general for any value a (that does not necessarily have anything to do with Y)., we use an indicator function ($I_a(\theta) = 1$ if $a = \theta$; 0 otherwise) because each latent value has (hypothetically) been observed exactly once together with the observed data point x_1. The probability of observing the latent value together with x_1 follows the posterior distribution in (6.10). This distribution is exactly the posterior in (6.10). For instance, $\mathbb{E}[I_{0.4}(\theta)] = P(\Theta = 0.4|X_1 = 6)$ and likewise for the other two mixture components.

We compute these expectations for each data point and add them up. We can do this because each latent variable is associated with exactly one observed variable, i.e. the observed data point x_i and is associated with latent data point y_i. Furthermore, we are computing the expectation of count statistics and counts are additive. This means that we can break down the expected counts over the entire data set into a sum of counts over data points. Notice that we are only able to do this because the statistic that we are computing decomposes nicely in accordance with the i.i.d. assumption over our data. If we computed the expectation of some other statistic,

say the maximum number of times that a coin has been used to generate consecutive sequences, this would not be possible.

The sum over expected counts for each data point gives us the expectations over the whole data set x and completes the E-step. In the M-step we assume that these expected values are the actual counts of how often we have observed each latent value. Let us call the counts $c_{0.4}, c_{0.5}, c_{0.65}$ where the index points to the corresponding mixture component. According to our model, the mixture components are multinomially distributed and thus in the M-step we want to find the MLE of that multinomial. In general, the MLE for θ_i of a multinomial is $\frac{c_i}{n}$[1]. In our case $n = 20$. Thus we set $\theta_i^{(1)} = \frac{c_i}{n}$ and complete the M-step. With our new parameter estimates, we can proceed to the second iteration of EM.

6.3 Basics of Information Theory

When we talk about *information*, we often use the term in qualitative sense. We say things like *This is valuable information* or *We have a lack of information*. We can also make statements about some information being more helpful than other. For a long time, however, people have been unable to quantify information. The person who succeeded in this endeavour was Claude E. Shannon who with his famous 1948 article *A Mathematical Theory of Communication* single-handedly created a new discipline: Information Theory! He also revolutionised digital communication and can be seen as one of the main contributors to our modern communication systems like the telephone, the internet etc.

The beauty about information theory is that it is based on probability theory and many results from probability theory seamlessly carry over to information theory. In this chapter, we are going to discuss the bare basics of information theory. These basics are often enough to understand many information theoretic arguments that researchers make in fields like computer science, psychology and linguistics.

Shannon's idea of information is as simple as it is compelling. Intuitively, if we are observing a realisation of a random variable, this realisation is surprising if it is unlikely to occur according to the distribution of that random variable. However, if the probability for the realisation is very low, than on average it does not occur very often, meaning that if we sample from the RV repeatedly, we are not surprised very often. We are not surprised when the probability mass of the distribution is concentrated on only a small subset of its support.

On the other hand, we quite often are surprised, if we cannot predict

---

[1]This fact can easily be seen by letting c_i be the number of successes in the realisation of a binomial RV and the sum of all c_j, $j \neq i$ be the number of failures. Then clearly $\frac{c_i}{n}$ is the MLE.

what the outcome of our next draw from the RV might be. We are surprised when the distribution over values of the RV is (close to) uniform. Thus, we are going to be most surprised on average if we are observing realisations of a uniformly distributed RV.

Shannon's idea was that observing RVs that cause a lot of surprises is informative because we cannot predict the outcomes and with each new outcome we have effectively learned something (namely that the i^{th} outcome took on the value that it did). Observing RVs with very concentrated distributions is not very informative under this conception because by just choosing the most probable outcome we can correctly predict most actually observed outcomes. Obviously, if I manage to predict an outcome beforehand, it's occurrence is not teaching me anything.

The goal of Shannon was to find a function that captures this intuitive idea. He eventually found it and showed that it is the only function to have properties that encompass the intuition. This function is called the **entropy** of a RV and it is simply the expected **surprisal** value.

Definition 6.4 (Surprisal) *The surprisal (value) of an outcome $x \in$ supp(X) of some RV X is defined as $-\log_2(P(X = x))$.*

Notice that we are using the logarithm of base 2 here. This is because surprisal and entropy are standardly measured in bits. Intuitively, the surprisal measures how many bits one needs to encode an observed outcome given that one knows the distribution underlying that outcome. The entropy measures how many bits one will need on average to encode an outcome that is generated by the distribution P_X.

Definition 6.5 (Entropy) *The entropy $H(P_X)$ of a RV X with distribution P_X is defined as*

$$H(P_X) := \mathbb{E}[-\log_2(P(X = x))] = - \sum_{x \in \text{supp}(X)} P(X = x) \log_2(P(X = x)) .$$

For the ease of notation, we often write $H(X)$ instead of $H(P_X)$.

Figure 6.1 shows the entropy of the Bernoulli distribution as a function of the parameter θ. The entropy function of the Bernoulli is often called the **binary entropy**. It measures the information of a binary decision, like a coin flip or an answer to a yes/no-question. The entropy of the Bernoulli is 1 bit when the distribution is uniform, i.e. when both choices are equally probable.

From the plot is it also easy to see that entropy is never negative. It holds in general that entropy is non-negative, because entropy is defined as expectation of surprisal and surprisal is the negative logarithm of probabilities. Because $\log(x) \leq 0$ for $x \in (0, 1]$, it is clear that $-\log(x) \geq 0$ for x in

Figure 6.1: Binary entropy function.

the same interval. Notice that from here on we drop the subscript and by convention let $\log = \log_2$.

A standard interpretation of the entropy is that it quantifies uncertainty. As we have pointed out before, a uniform distribution means that you are most uncertain and indeed the uniform distribution maximizes the entropy. However, the more choices you have to pick from, the more uncertain you are going to be. The entropy function also captures this intuition. Notice that if a discrete distribution is uniform, all probabilities are $\frac{1}{|\operatorname{supp}(X)|}$. Clearly, as we increase $|\operatorname{supp}(X)|$, we decrease the probabilities. By decreasing the probabilities, we increase their negative logarithms, and hence their surprisal. Let us make this intuition more formal.

Theorem 6.6 *A discrete RV X with uniform distribution and support of size n has entropy $H(X) = \log(n)$.*

Proof:

$$(6.11) \qquad H(X) = \sum_{x \in \operatorname{supp}(X)} -\log(P(X = x))P(X = x)$$

$$(6.12) \qquad \qquad = \sum_{x \in \operatorname{supp}(X)} \log(n)P(X = x) = \log(n) \,. \quad \square$$

9

Exercise 6.7 *You are trying to learn chess and you start by studying where chess grandmasters move their king when it is positioned in one of the middle fields of the board. The king can move to any of the adjoining 8 fields. Since you do not know a thing about chess yet, you assume that each move is equally probable. In this situation, what is the entropy of moving the king?*

At the outset of this section we promised you that you could easily transfer results from probability theory to information theory. We will not be able to show any kind of linearity for entropy because it contains log-terms and the logarithm is not linear. We can however find alternative expressions for joint entropy (where the joint entropy is simply the entropy of a joint RV). Before we do so, let us also define the notion of conditional entropy. We have seen in Section **??** that $P_{X|Y=y}$ is a valid probability distribution for any $y \in \text{supp}(Y)$ such that $P(Y = y) > 0$. Hence, we can also define its conditional entropy.

Definition 6.8 (Conditional Entropy) *For two jointly distributed RVs X, Y and $y \in \text{supp}(Y)$ such that $P(Y = y) > 0$, the conditional entropy of X given that $Y = y$ is defined as*

$$
\begin{aligned}
H(X|Y = y) &:= \mathbb{E}_X[-\log_2(P(X = x|Y = y))] \\
&= -\sum_{x \in \text{supp}(X)} P(X = x|Y = y) \log_2(P(X = x|Y = y)).
\end{aligned}
$$

The conditional entropy of X given Y is defined as

$$
H(X|Y) := \mathbb{E}_Y[H(X|Y)] = \sum_{y \in \text{supp}(Y)} P(Y = y)H(X|Y = y).
$$

With this definition at hand we show that the joint entropy decomposes according to the chain rule.

$$H(X,Y) = \sum_{\substack{x\in\text{supp}(X)\\y\in\text{supp}(Y)}} -\log(P(X=x,Y=y)) \times P(X=x,Y=y)$$

$$= \sum_{\substack{x\in\text{supp}(X)\\y\in\text{supp}(Y)}} -\log(P(X=x|Y=y)) \times P(X=x,Y=y)$$

$$- \sum_{y\in\text{supp}(Y)} \log(P(Y=y)) \times \sum_{x\in\text{supp}(X)} P(X=x,Y=y)$$

$$= \sum_{y\in\text{supp}(Y)} P(Y=y) \times \sum_{x\in\text{supp}(X)} -\log(P(X=x|Y=y)) \times P(X=x|Y=y)$$

$$- \sum_{y\in\text{supp}(Y)} \log(P(Y=y)) \times P(Y=y)$$

$$= H(X|Y) + H(Y)$$

Exercise 6.9 *Prove that $H(X,Y|Z) = H(X|Z) + H(Y|Z)$ if $X \perp Y|Z$.*

Now that we have seen some information-theoretic concepts, you may be happy to hear that there is an information-theoretic interpretation of EM. This interpretation helps us to get a better intuition for the algorithm. To formulate that interpretation we need one more concept, however.

Definition 6.10 (Relative Entropy) *The relative entropy of RVs X, Y with distributions P_X, P_Y and $\text{supp}(X) \subseteq \text{supp}(Y)$ is defined as*

$$D(P_X||P_Y) := \sum_{x\in\text{supp}(X)} P(X=x) \log \frac{P(X=x)}{P(Y=x)} \ .$$

If $P(X=y) = 0$ for any $y \in \text{supp}(Y)$ we define $D(P_X||P_Y) = \infty$. As with entropy, we often abbreviate $D(P_X||P_Y)$ with $D(X||Y)$.

The relative entropy is commonly known as **Kullback-Leibler (KL)** divergence. It measures the entropy of X as scaled to Y. Intuitively, it gives a measure of how "far away" P_X is from P_Y. To understand "far away", recall that entropy is a measure of uncertainty. The relative entropy measure the uncertainty that you have about P_X if you know P_Y **Chris:** hard to see why at this point. This uncertainty is low if both distributions place most of their mass on the same outcomes. Since $\log(1) = 0$ the relative entropy is 0 if $P_X = P_Y$.

It is worthwhile to point out the difference between relative and conditional entropy. Conditional entropy is the average entropy of X given that you know what value Y takes on. In the case of relative entropy you do not know the value of Y, only its distribution.

> **Exercise 6.11** *Show that $D(X, Y || Y) = H(X|Y)$. Furthermore show that $D(X, Y || Y) = H(X)$ if $X \perp Y$.*

Further Material

At the ILLC, the best place to learn more about information theory is Christan Schaffner's course that is taught every year. David MacKay also offers a free book on the subject. Finally, Coursera also offers an online course on information theory.

To get a better understanding of EM and the other concepts discussed in this script along with some more examples, consult Micheal Collins' lecture notes .